


Generative models

Outline

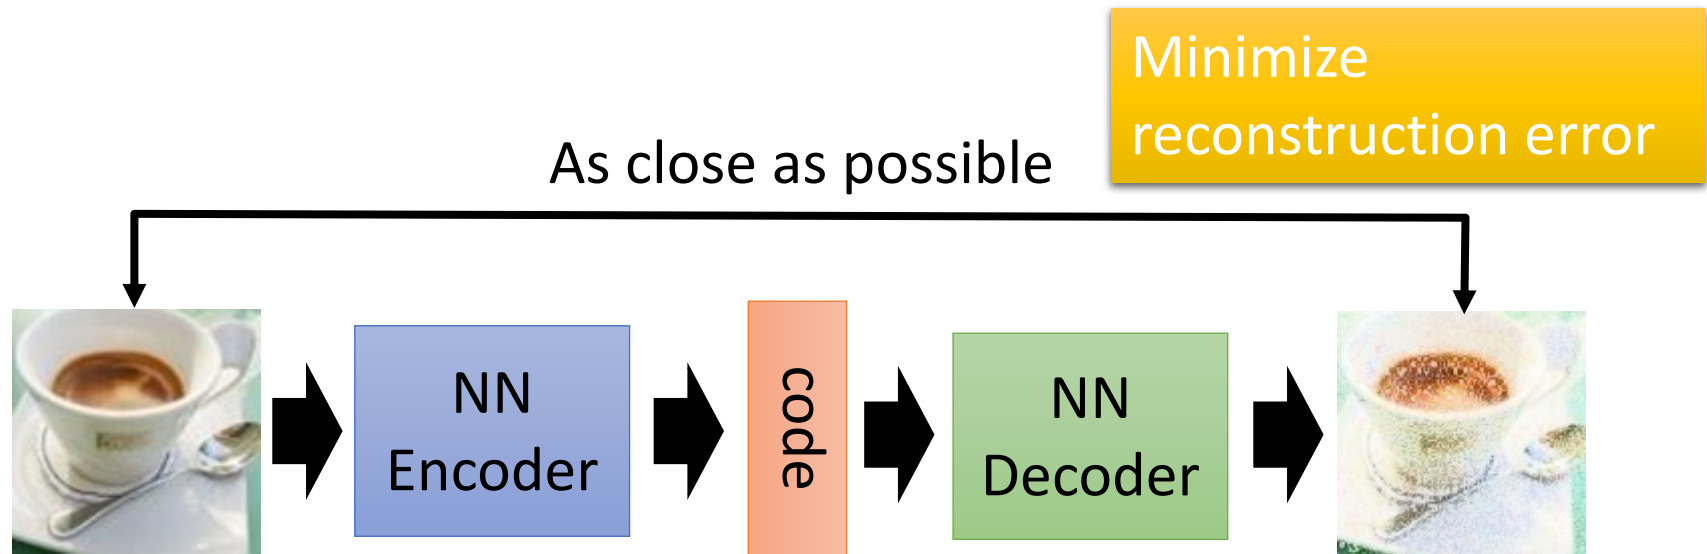
1. **Preview: Auto-Encoders, VAE**
2. Generative models with GAN
3. GAN architectures




Drawing? => learning from examples



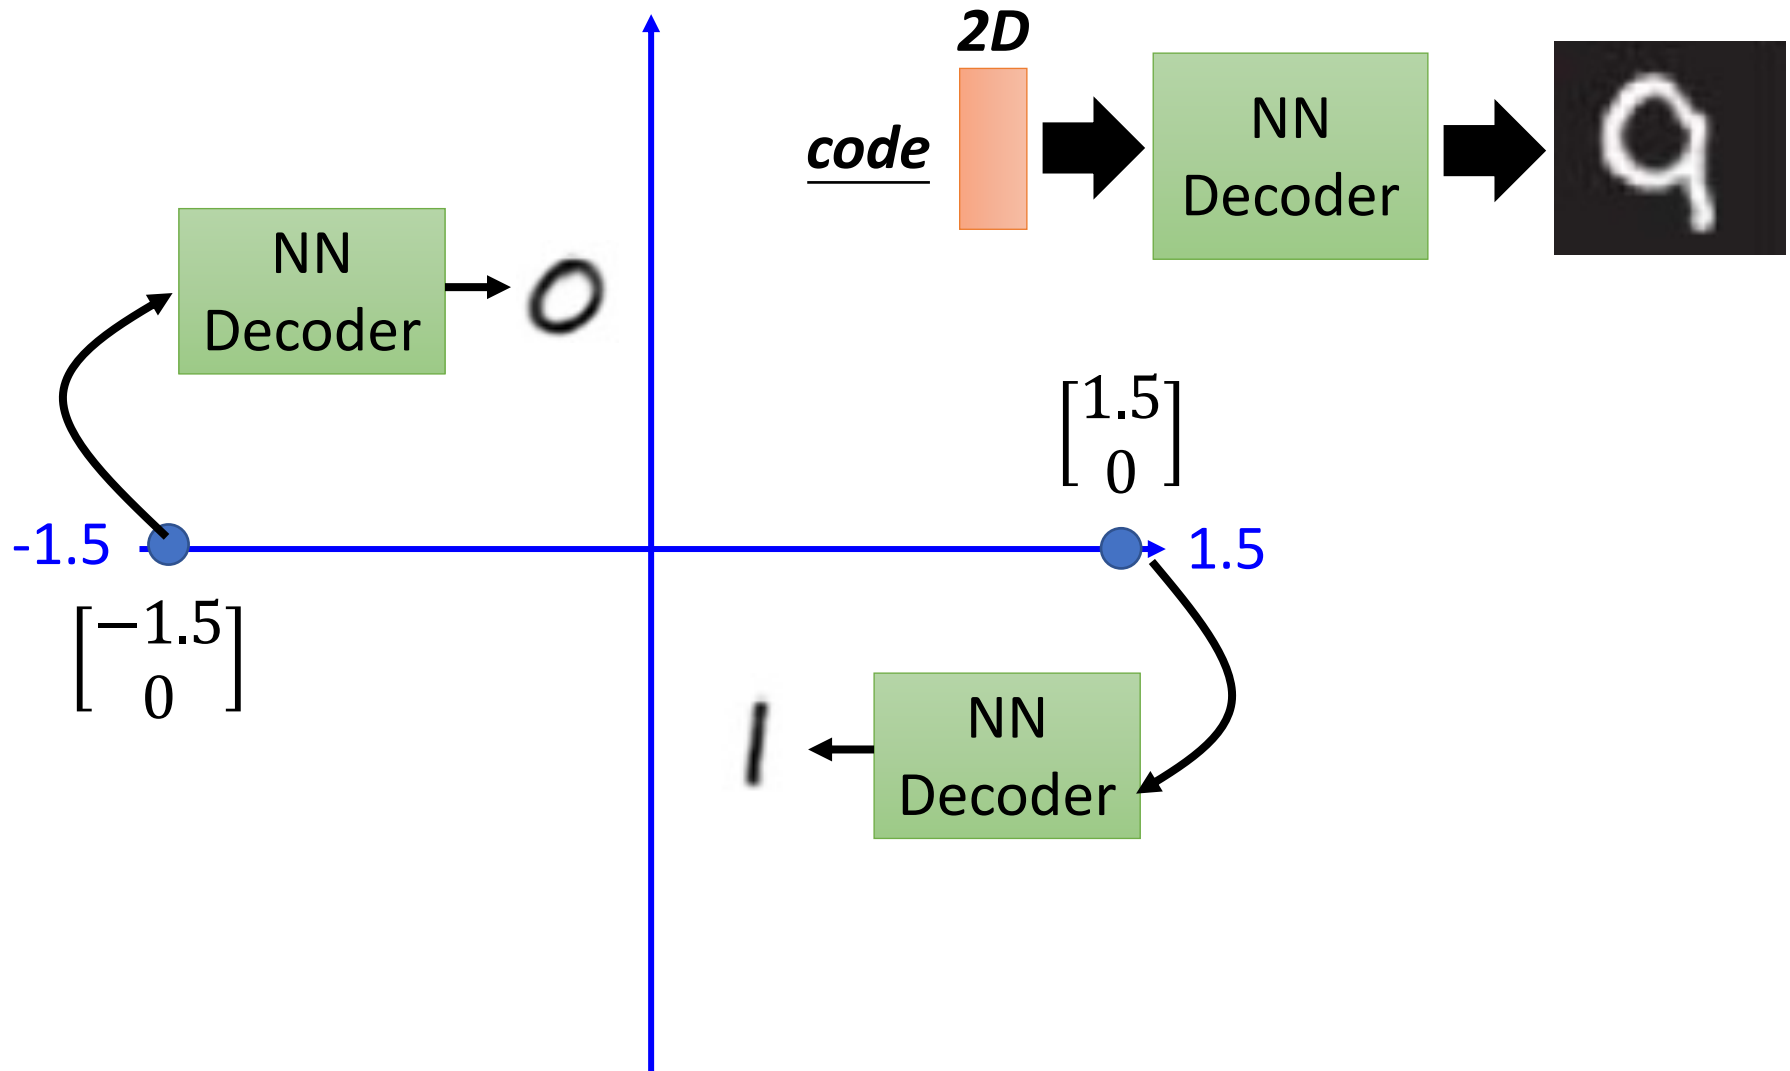
Review: Auto-encoder




Randomly generate
a vector as code




Review: Auto-encoder



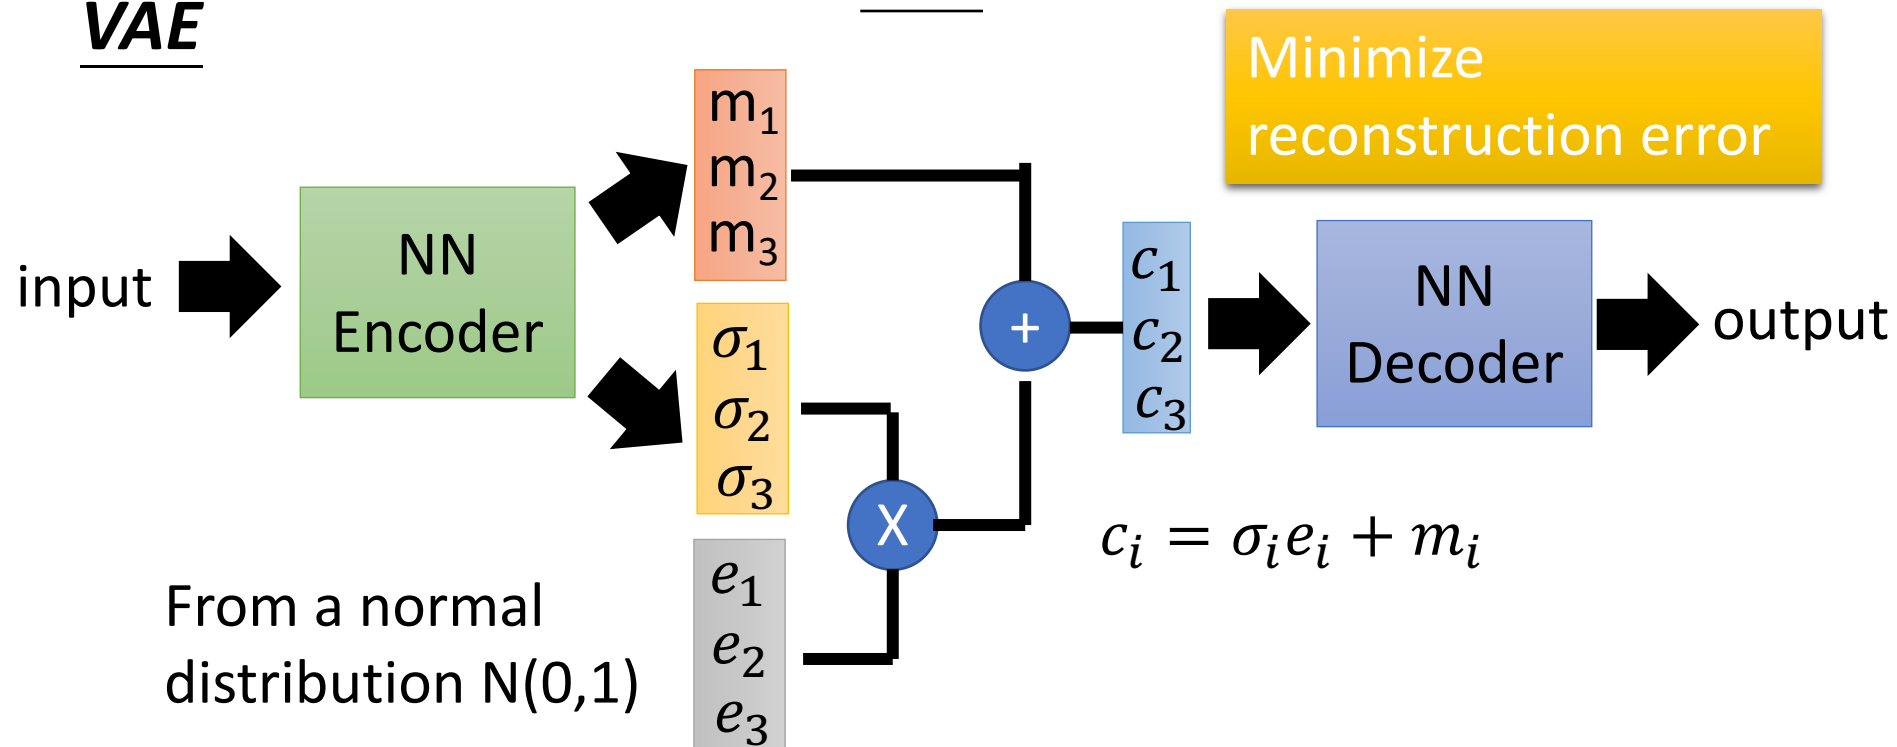
Review: Auto-encoder



Auto-encoder

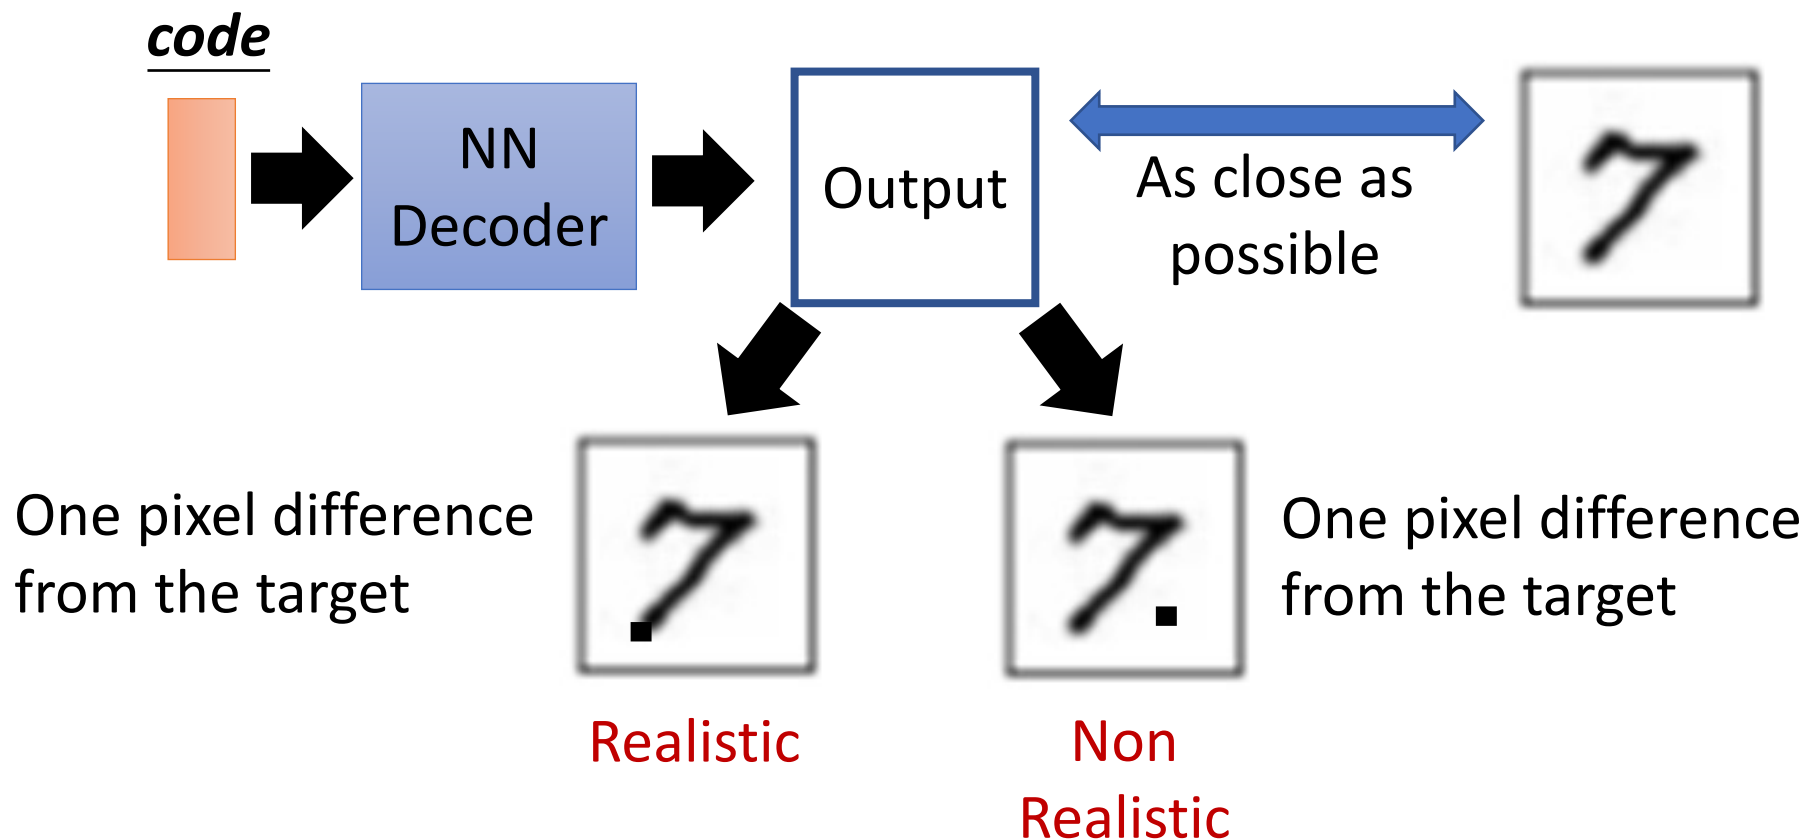


VAE




Problems of AE/VAE

- It does not really try to simulate real images




Problems of AE/VAE

GAN to tackle this pb:



Realistic



Non Realistic

GAN: generative **adversarial** networks

Game scenario:

Player1, Generator, produces samples

Player2, – Its adversary **Discriminator**, attempts to distinguish **real** samples from **fake** generated ones (produced by Player1) !

Player1 aims at producing **Realistic** images to fool the Player2

Generative models

Outline

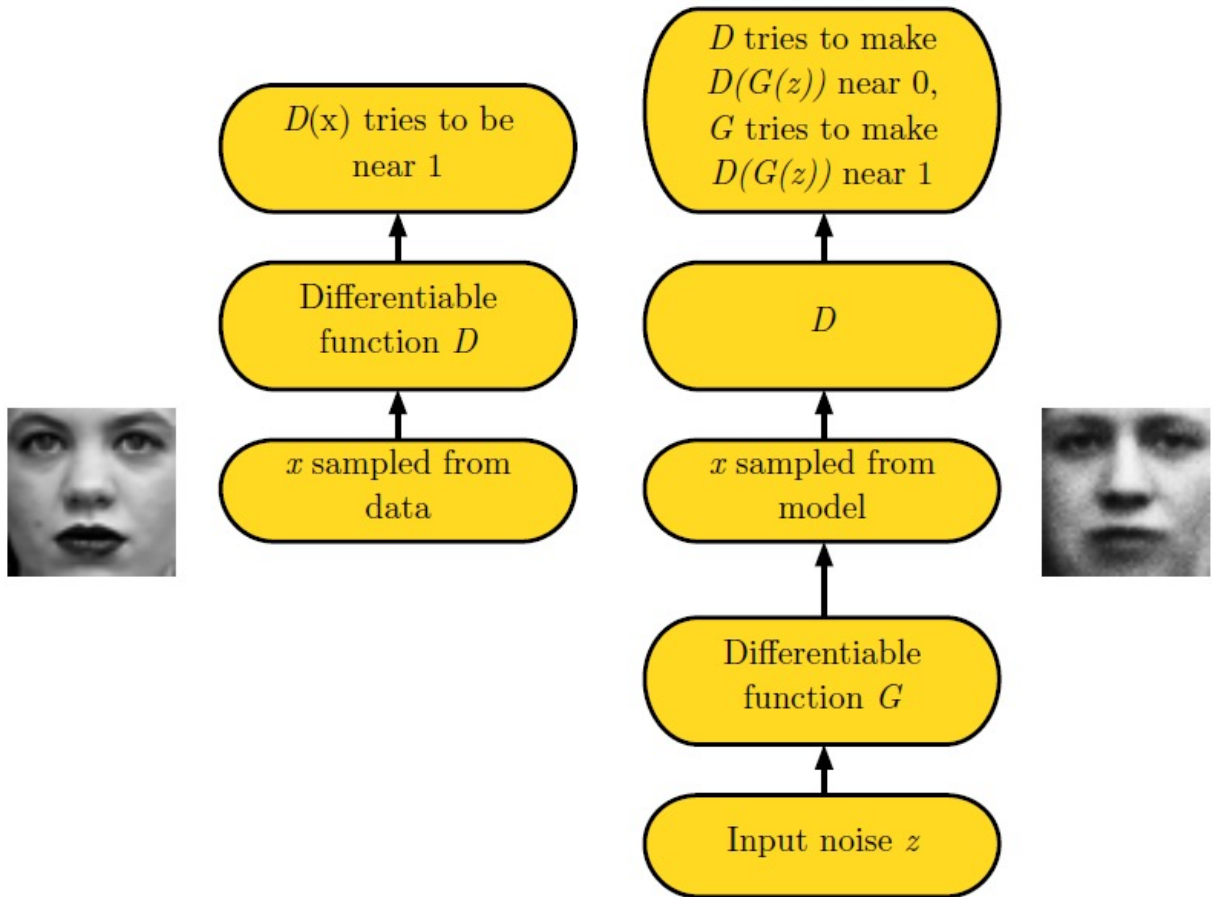
1. Preview: Auto-Encoders, VAE
- 2. Generative models with GAN**
 - GAN Algorithm

Adversarial Nets Framework

Game scenario:

Player1, **Generator G**

Player2, Discriminator D



$$V(G, D) = \mathbb{E}_{x \sim P_{data}} [\log D(x)] + \mathbb{E}_{x \sim P_G} [\log(1 - D(x))]$$

$$G^* = \arg \min_G \max_D V(G, D)$$

GAN Learning – D and G updates

Game scenario:

Player1, Generator G, produces samples

Player2, – Its adversary **Discriminator D**, attempts to distinguish **real** samples from **fake** generated ones (produced by Player1) !

Player1 aims at producing **Realistic** images to fool the Player2


NN
Generator
v1



Discri-
minator
v1

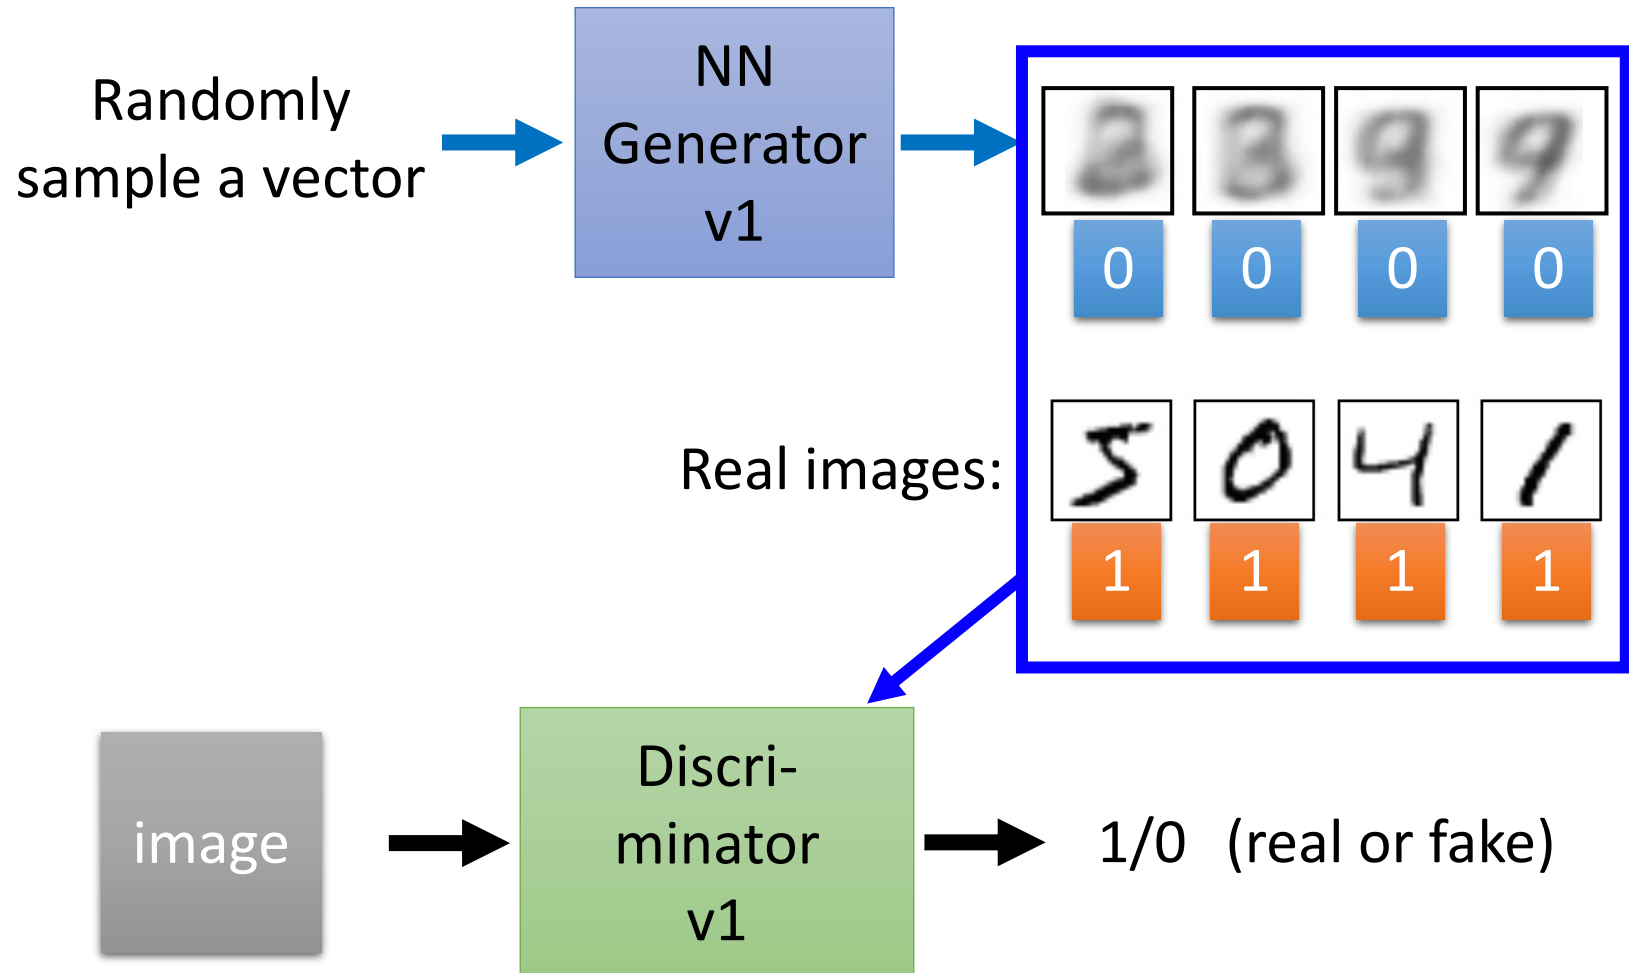
Binary
Classifier

Real images:



Fake images:

GAN - Discriminator



Discriminator Optimization on a batch of images:

Using gradient descent to update the parameters in the discriminator, with a fixed generator

GAN - Generator

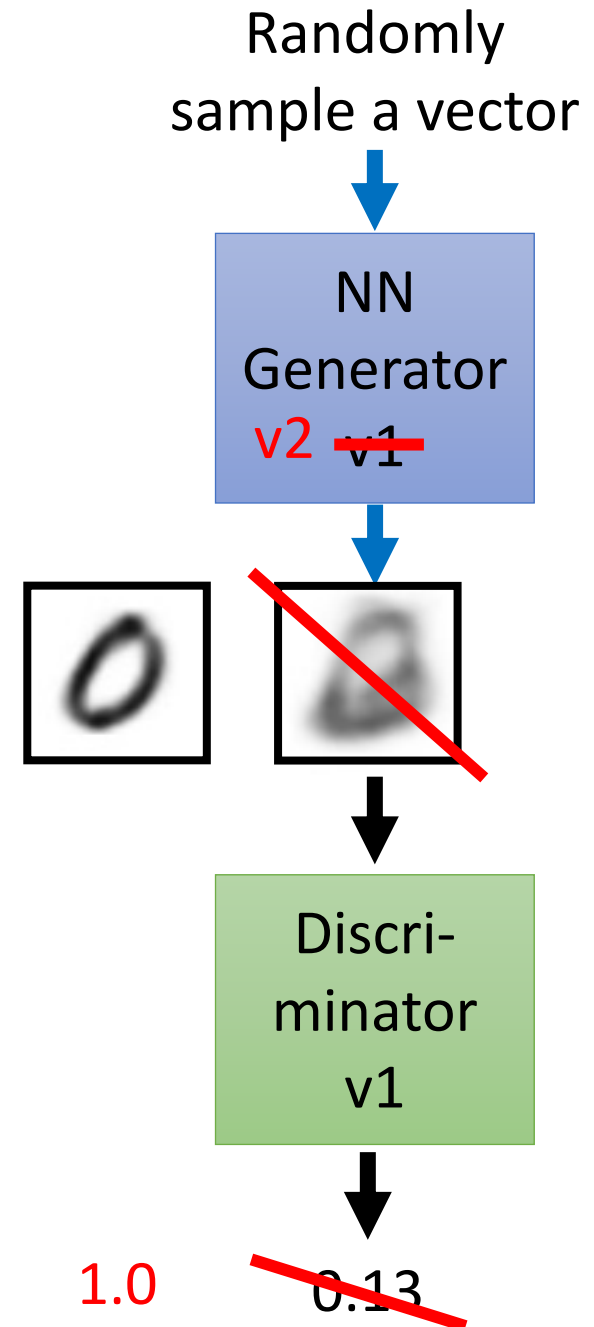
Updating the parameters of generator

➡ The output be classified as “real” (as close to 1 as possible)

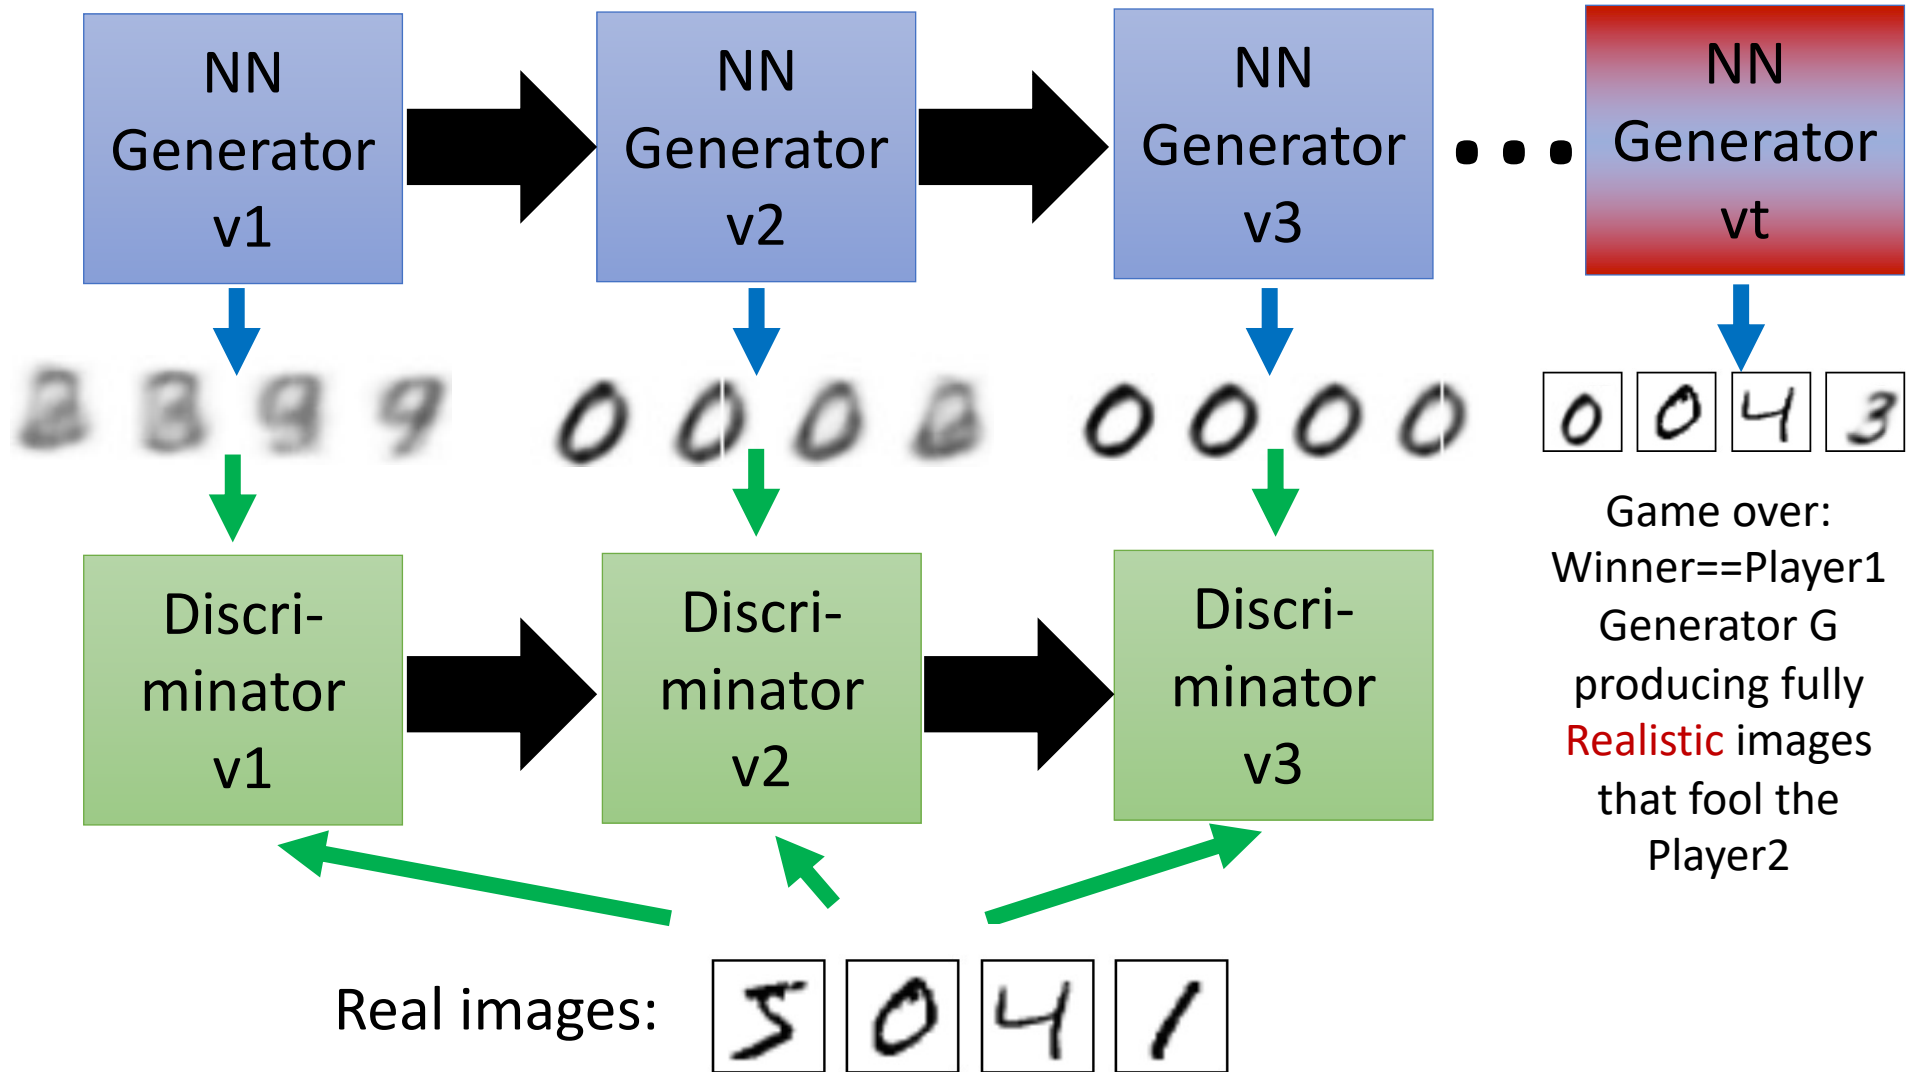
Generator + Discriminator
= a network

Optimization:

Using gradient descent to update the parameters in the generator, but fixing the discriminator



GAN Learning – D and G updates



Algorithm 1 Minibatch stochastic gradient descent training of generative adversarial nets. The number of steps to apply to the discriminator, k , is a hyperparameter. We used $k = 1$, the least expensive option, in our experiments.

for number of training iterations **do**

for k steps **do**

- Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$.
- Sample minibatch of m examples $\{x^{(1)}, \dots, x^{(m)}\}$ from data generating distribution $p_{\text{data}}(x)$.
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[\log D(x^{(i)}) + \log (1 - D(G(z^{(i)}))) \right].$$

end for

- Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$.
- Update the generator by descending its stochastic gradient:

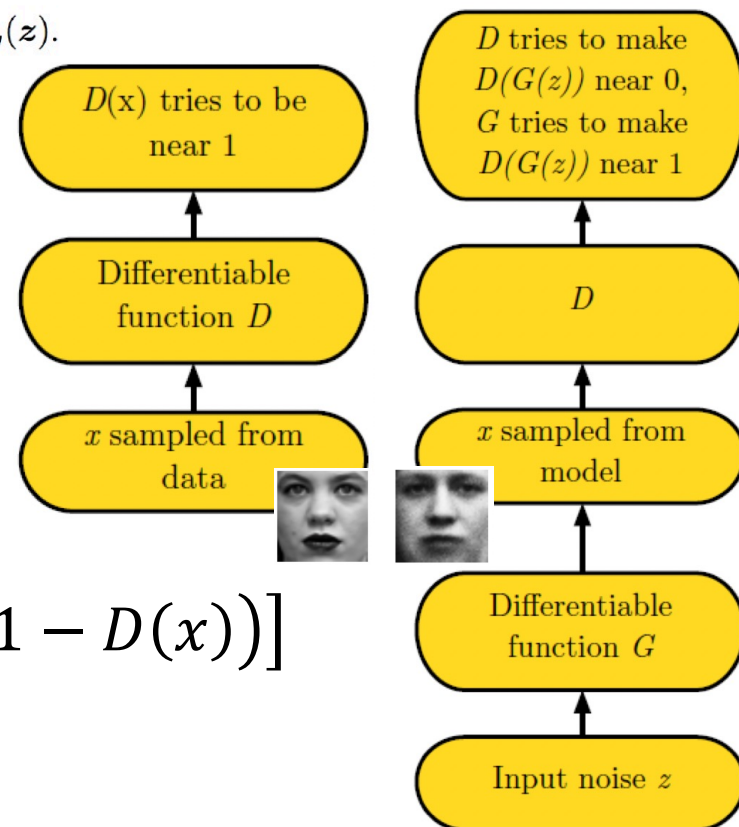
$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (1 - D(G(z^{(i)}))).$$

end for

GAN algorithm

$$V = \mathbb{E}_{x \sim P_{\text{data}}} [\log D(x)] + \mathbb{E}_{x \sim P_G} [\log (1 - D(x))]$$

$$G^* = \arg \min_G \max_D V(G, D)$$



One example GAN

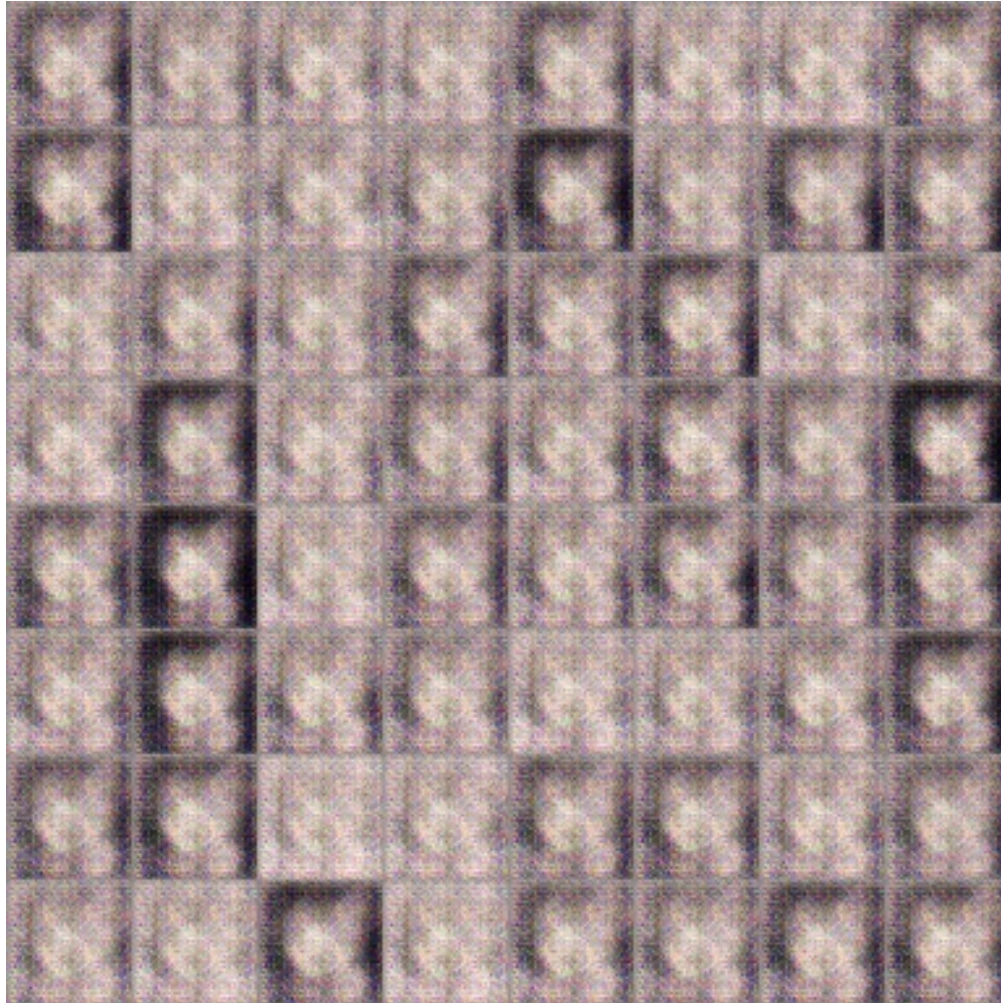


Source of images: <https://zhuanlan.zhihu.com/p/24767059>

DCGAN: <https://github.com/carpedm20/DCGAN-tensorflow>

GAN

100 rounds



GAN



1000 rounds

GAN



2000 rounds

GAN




5000 rounds

GAN



10,000 rounds

GAN



20,000 rounds

GAN



50,000 rounds

Generative models

Outline

1. Preview: Auto-Encoders, VAE
2. Generative models with GAN
 - GAN Algorithm
 - **KL vs. Jensen Shanon Divergence**

$$V(\textcolor{red}{G}, D) = \mathbb{E}_{x \sim P_{data}} [\log D(x)] + \mathbb{E}_{x \sim P_{\textcolor{red}{G}}} [\log(1 - D(x))]$$

$$\textcolor{red}{G}^* = \arg \min_{\textcolor{red}{G}} \max_D V(\textcolor{red}{G}, D)$$

Which measure to evaluate how $P_G(x; \theta)$ is close to $P_{data}(x)$ in Maximum Likelihood optimization?

- Given a data distribution $P_{data}(x)$
- We have a distribution $P_G(x; \theta)$ parameterized by θ
 - E.g. $P_G(x; \theta)$ is a Gaussian Mixture Model, θ are means and variances of the Gaussians
 - We want to find θ such that $P_G(x; \theta)$ close to $P_{data}(x)$

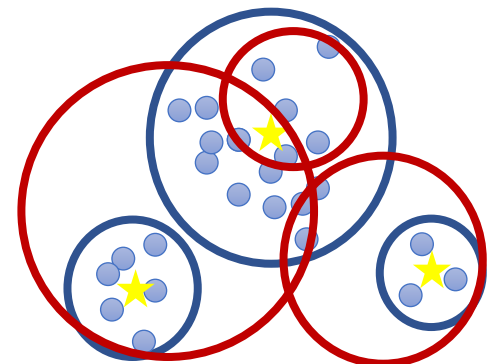
Sample $\{x^1, x^2, \dots, x^m\}$ from $P_{data}(x)$

We can compute $P_G(x^i; \theta)$

Likelihood of generating the samples

$$L = \prod_{i=1}^m P_G(x^i; \theta)$$

Find θ^* maximizing the likelihood



Which measure to evaluate how $P_G(x; \theta)$ is close to $P_{data}(x)$ in Maximum Likelihood optimization?

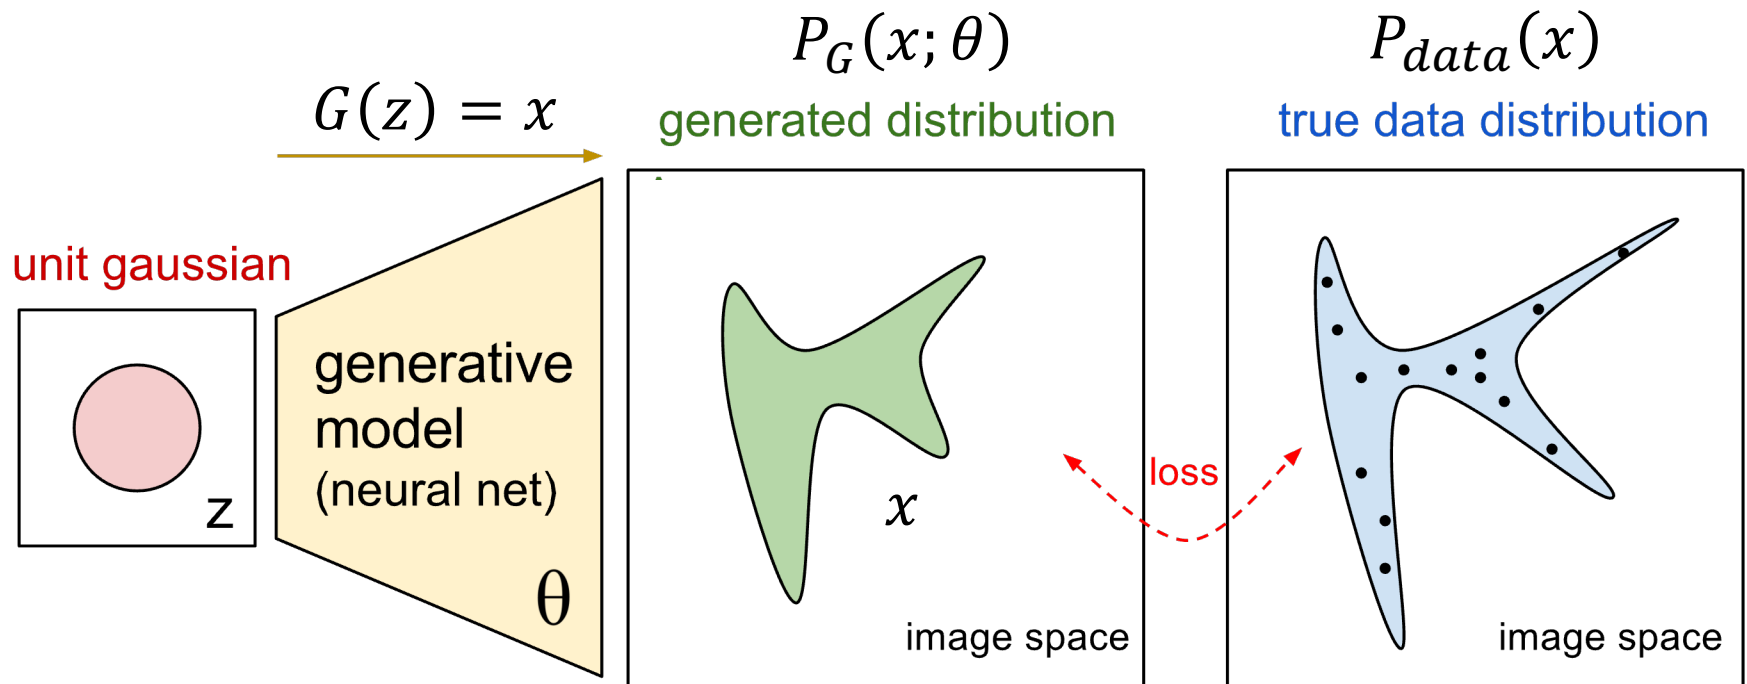
$$\begin{aligned}\theta^* &= \arg \max_{\theta} \prod_{i=1}^m P_G(x^i; \theta) = \arg \max_{\theta} \log \prod_{i=1}^m P_G(x^i; \theta) \\ &= \arg \max_{\theta} \sum_{i=1}^m \log P_G(x^i; \theta) \quad \{x^1, x^2, \dots, x^m\} \text{ from } P_{data}(x)\end{aligned}$$

$$\approx \arg \max_{\theta} \mathbb{E}_{x \sim P_{data}} [\log P_G(x; \theta)]$$

$$\begin{aligned}&= \arg \max_{\theta} \int_x P_{data}(x) \log P_G(x; \theta) dx - \int_x P_{data}(x) \log P_{data}(x) dx \\ &= \arg \min_{\theta} KL(P_{data}(x) || P_G(x; \theta)) \quad KL(P || Q) = \int_x P(x) \log \frac{P(x)}{Q(x)} dx\end{aligned}$$

In Maximum Likelihood it is a KLD Kullback Leibler Divergence

If $P_G(x; \theta)$ is a coming with a NN



$$P_G(x; \theta) = \int_z P_{prior}(z) I_{[G(z)=x]} dz$$

It is difficult to compute the likelihood.

Basic Idea of GAN: the 2 players G-D game

- Generator G Hard to learn by maximum likelihood
 - G is a function, input z , output x
 - Given a prior distribution $P_{\text{prior}}(z)$, a probability distribution $P_G(x)$ is defined by function G (and P_{prior})
- Discriminator D
 - D is a function, input x , output scalar
 - Evaluate the “difference” between $P_G(x)$ and $P_{\text{data}}(x)$
- Global objective function $V(G, D)$

$$\theta^* = G^* = \arg \min_G \max_D V(G, D)$$


Basic Idea

$$\boxed{G^*} = \arg \min_G \max_D V(G, D)$$

$$V = \mathbb{E}_{x \sim P_{data}} [\log D(x)] + \mathbb{E}_{x \sim P_G} [\log(1 - D(x))]$$

Given a generator G , $\max_D V(G, D)$ evaluate the “difference” between P_G and P_{data}

Pick the G defining P_G most similar to P_{data}



$$\max_D V(G, D) \quad G^* = \arg \min_G \max_D V(G, D)$$

- Given G , what is the optimal D^* maximizing

$$\begin{aligned} V &= \mathbb{E}_{x \sim P_{data}} [\log D(x)] + \mathbb{E}_{x \sim P_G} [\log(1 - D(x))] \\ &= \int_x P_{data}(x) \log D(x) dx + \int_x P_G(x) \log(1 - D(x)) dx \\ &= \int_x [P_{data}(x) \log D(x) + P_G(x) \log(1 - D(x))] dx \end{aligned}$$

Assume that $D(x)$ can have any value here

- Given x , the optimal D^* maximizing

$$P_{data}(x) \log D(x) + P_G(x) \log(1 - D(x))$$

$$\max_D V(G, D) \quad G^* = \arg \min_G \max_D V(G, D)$$

- Given x , the optimal D^* maximizing

$$\underbrace{P_{data}(x)}_a \log \underbrace{D(x)}_D + \underbrace{P_G(x)}_b \log \underbrace{(1 - D(x))}_D$$

- Find D^* maximizing: $f(D) = a \log(D) + b \log(1 - D)$

$$\frac{df(D)}{dD} = a \times \frac{1}{D} + b \times \frac{1}{1 - D} \times (-1) = 0$$

$$a \times \frac{1}{D^*} = b \times \frac{1}{1 - D^*}$$

$$a \times (1 - D^*) = b \times D^*$$

$$a - aD^* = bD^*$$


$$D^* = \frac{a}{a + b} \rightarrow$$

$$D^*(x) = \frac{P_{data}(x)}{P_{data}(x) + P_G(x)} \quad 0 < \quad < 1$$

$$\max_D V(G, D) \quad G^* = \arg \min_G \max_D V(G, D)$$

$$D_1^*(x) = \frac{P_{data}(x)}{P_{data}(x) + P_{G_1}(x)}$$

$$D_2^*(x) = \frac{P_{data}(x)}{P_{data}(x) + P_{G_2}(x)}$$



$$\max_D V(G, D)$$

$$V = \mathbb{E}_{x \sim P_{data}} [\log D(x)] + \mathbb{E}_{x \sim P_G} [\log(1 - D(x))]$$

$$\max_D V(G, D) = V(G, D^*)$$

$$D^*(x) = \frac{P_{data}(x)}{P_{data}(x) + P_G(x)}$$

$$= \mathbb{E}_{x \sim P_{data}} \left[\log \frac{P_{data}(x)}{P_{data}(x) + P_G(x)} \right] + \mathbb{E}_{x \sim P_G} \left[\log \frac{P_G(x)}{P_{data}(x) + P_G(x)} \right]$$

$$= \int_x P_{data}(x) \log \frac{\frac{1}{2} P_{data}(x)}{\frac{P_{data}(x) + P_G(x)}{2}} dx + \int_x P_G(x) \log \frac{\frac{1}{2} P_G(x)}{\frac{P_{data}(x) + P_G(x)}{2}} dx$$

$$\Rightarrow +2 \log \frac{1}{2} = -2 \log 2$$

$$\max_D V(G, D)$$

$$\begin{aligned} \text{JSD}(P||Q) &= \frac{1}{2}\text{KL}(P||M) + \frac{1}{2}\text{KL}(Q||M) \\ M &= \frac{1}{2}(P + Q) \end{aligned}$$

$$\max_D V(G, D) = V(G, D^*)$$

$$D^*(x) = \frac{P_{data}(x)}{P_{data}(x) + P_G(x)}$$

$$= -2\log 2 + \int_x P_{data}(x) \log \frac{P_{data}(x)}{(P_{data}(x) + P_G(x))/2} dx$$

$$+ \int_x P_G(x) \log \frac{P_G(x)}{(P_{data}(x) + P_G(x))/2} dx$$

$$\begin{aligned} = -2\log 2 + \text{KL} \left(P_{data}(x) || \frac{P_{data}(x) + P_G(x)}{2} \right) \\ + \text{KL} \left(P_G(x) || \frac{P_{data}(x) + P_G(x)}{2} \right) \end{aligned}$$

$$= -2\log 2 + 2\text{JSD}(P_{data}(x)||P_G(x)) \quad \text{Jensen-Shannon divergence}$$

In the end

$$V = \mathbb{E}_{x \sim P_{data}} [\log D(x)] \\ + \mathbb{E}_{x \sim P_G} [\log(1 - D(x))]$$

- Generator G, Discriminator D
- Looking for G^* such that $G^* = \arg \min_G \max_D V(G, D)$
- Given G, $\max_D V(G, D) = -2\log 2 + 2JSD(P_{data}(x) || P_G(x))$
 $0 < \hspace{15em} < \log 2$
- What is the optimal G?

$$P_G(x) = P_{data}(x)$$

with/using the $JS(P_G, P_{data})$ Divergence

(In Maximum Likelihood it is a KL Divergence)

